

Global Air Quality and Respiratory Health Outcomes

July 1, 2025

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[3]: df = pd.read_csv('../data/air_quality_health_dataset.csv')
print(df.head())
print(df.info())
```

	city	date	aqi	pm2_5	pm10	no2	o3	temperature \
0	Los Angeles	2020-01-01	65	34.0	52.7	2.2	38.5	33.5
1	Beijing	2020-01-02	137	33.7	31.5	36.7	27.5	-1.6
2	London	2020-01-03	266	43.0	59.6	30.4	57.3	36.4
3	Mexico City	2020-01-04	293	33.7	37.9	12.3	42.7	-1.0
4	Delhi	2020-01-05	493	50.3	34.8	31.2	35.6	33.5

	humidity	hospital_admissions	population_density	hospital_capacity
0	33	5	Rural	1337
1	32	4	Urban	1545
2	25	10	Suburban	1539
3	67	10	Urban	552
4	72	9	Suburban	1631

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 88489 entries, 0 to 88488

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	city	88489 non-null	object
1	date	88489 non-null	object
2	aqi	88489 non-null	int64
3	pm2_5	88489 non-null	float64
4	pm10	88489 non-null	float64
5	no2	88489 non-null	float64
6	o3	88489 non-null	float64
7	temperature	88489 non-null	float64
8	humidity	88489 non-null	int64
9	hospital_admissions	88489 non-null	int64
10	population_density	88489 non-null	object
11	hospital_capacity	88489 non-null	int64

```
dtypes: float64(5), int64(4), object(3)
memory usage: 8.1+ MB
None
```

```
[5]: # Missing values are checked for each column
print("Missing values:\n", df.isna().sum())
```

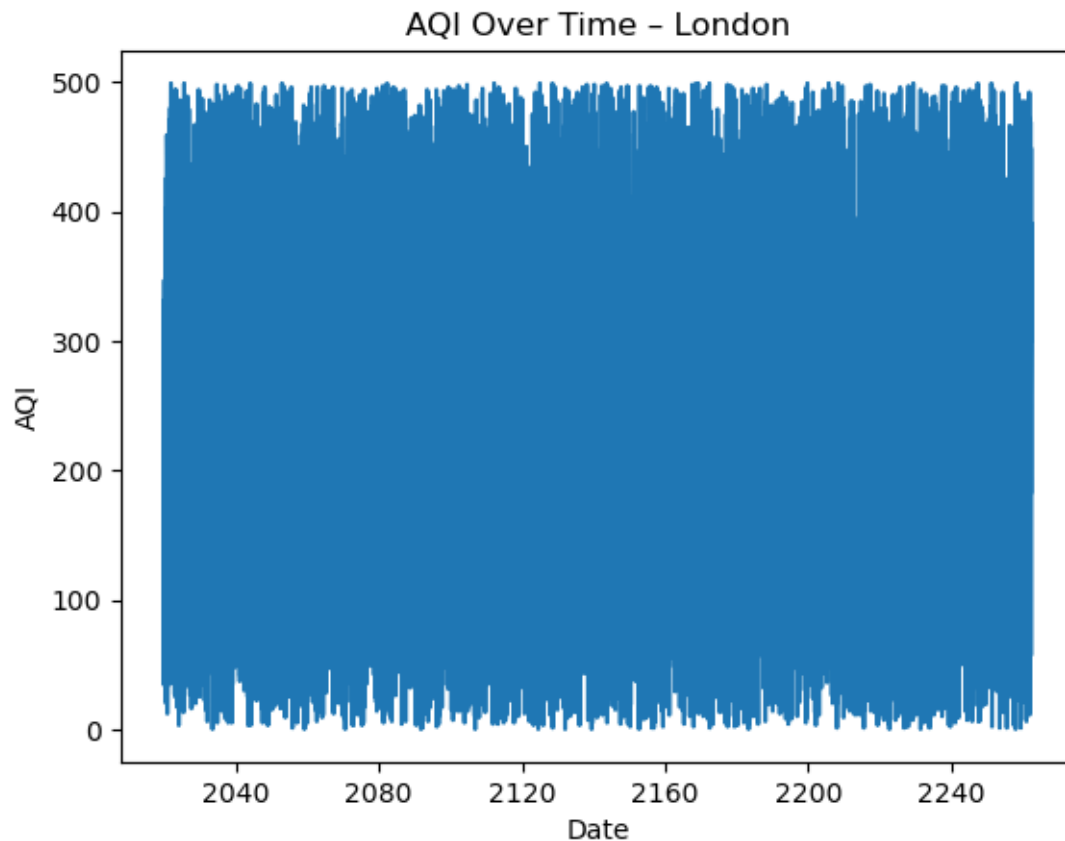
```
Missing values:
  city      0
  date      0
  aqi       0
  pm2_5     0
  pm10      0
  no2       0
  o3        0
  temperature  0
  humidity   0
  hospital_admissions  0
  population_density  0
  hospital_capacity  0
dtype: int64
```

```
[7]: # Convert the 'date' column to datetime format
df['date'] = pd.to_datetime(df['date'])
```

```
[15]: # Top 10 days with the highest PM2.5 levels
top_pm25 = df.nlargest(10, 'pm2_5')[['city', 'date', 'pm2_5']]
print(top_pm25)
```

	city	date	pm2_5
52628	Delhi	2164-02-03	109.9
43181	Delhi	2138-03-24	108.7
31482	Cairo	2106-03-13	105.2
5164	Mexico City	2034-02-20	103.8
75840	Delhi	2227-08-24	96.3
43350	Delhi	2138-09-09	95.9
10471	Mexico City	2048-09-01	94.7
40433	Delhi	2130-09-14	94.1
57637	Delhi	2177-10-21	93.4
77990	Delhi	2233-07-13	92.9

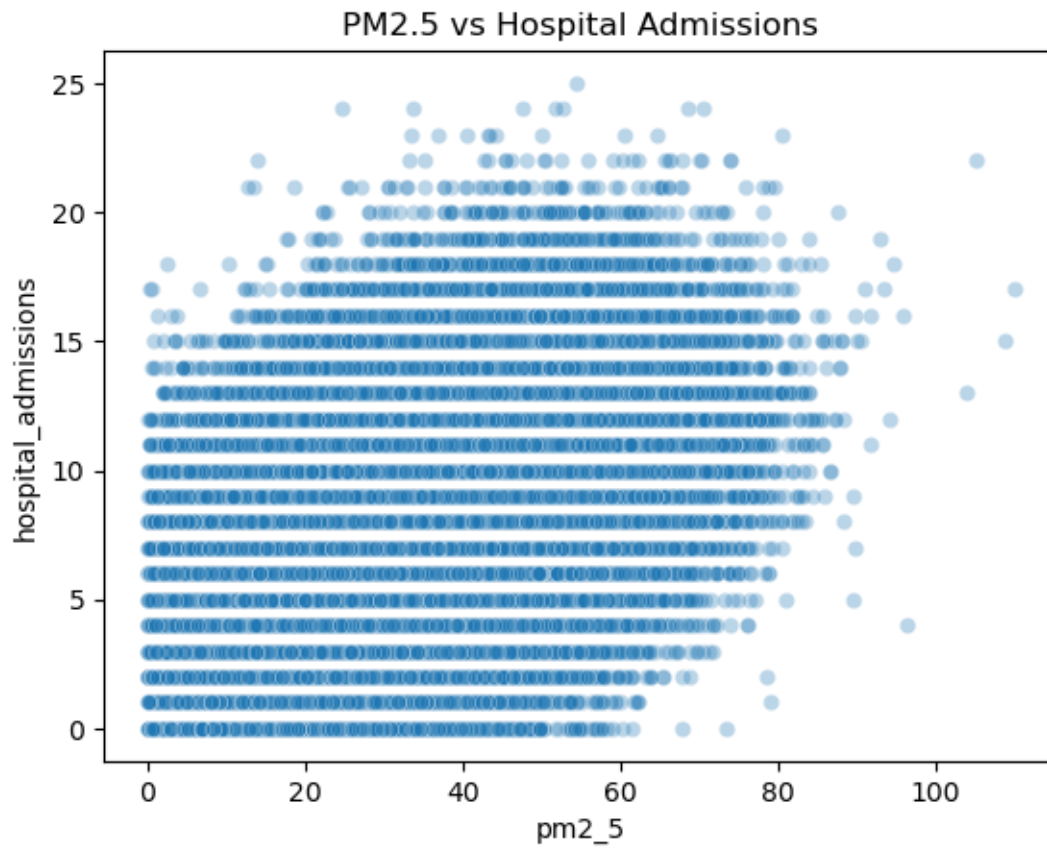
```
[17]: #Line plot of Air Quality Index (AQI) in London
la = df[df['city']=='London']
plt.plot(la['date'], la['aqi'])
plt.title('AQI Over Time - London')
plt.ylabel('AQI')
plt.xlabel('Date')
plt.show()
```



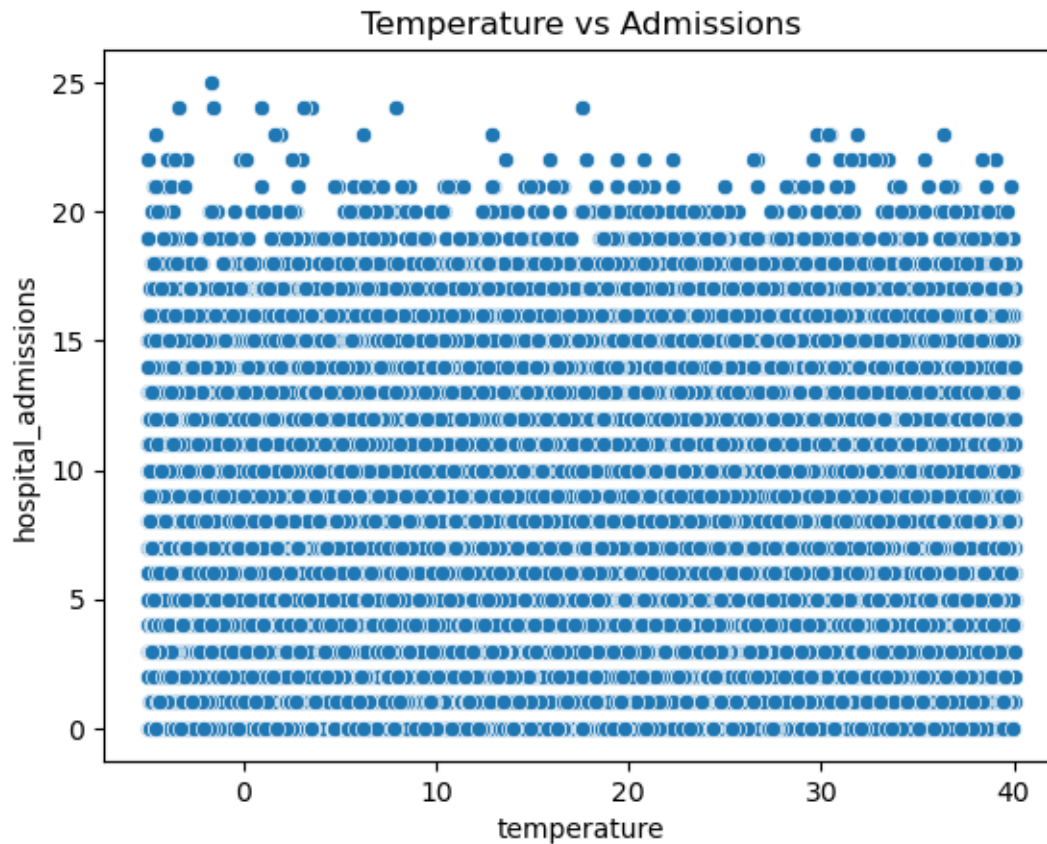
```
[35]: # Average hospital admissions per city (Top 10 cities)
df.groupby('city')['hospital_admissions'].mean().sort_values(ascending=False).
    ↪head(10)
```

```
[35]: city
Cairo      8.108519
Beijing    8.096039
London     8.057409
Delhi      8.044323
Tokyo      8.028469
Los Angeles 8.021215
Mexico City 8.007550
São Paulo  7.952517
Name: hospital_admissions, dtype: float64
```

```
[37]: # Scatter plot showing the relationship between PM2.5 levels and hospital_
      ↪admissions
      sns.scatterplot(data=df, x='pm2_5', y='hospital_admissions', alpha=0.3)
      plt.title('PM2.5 vs Hospital Admissions')
      plt.show()
```



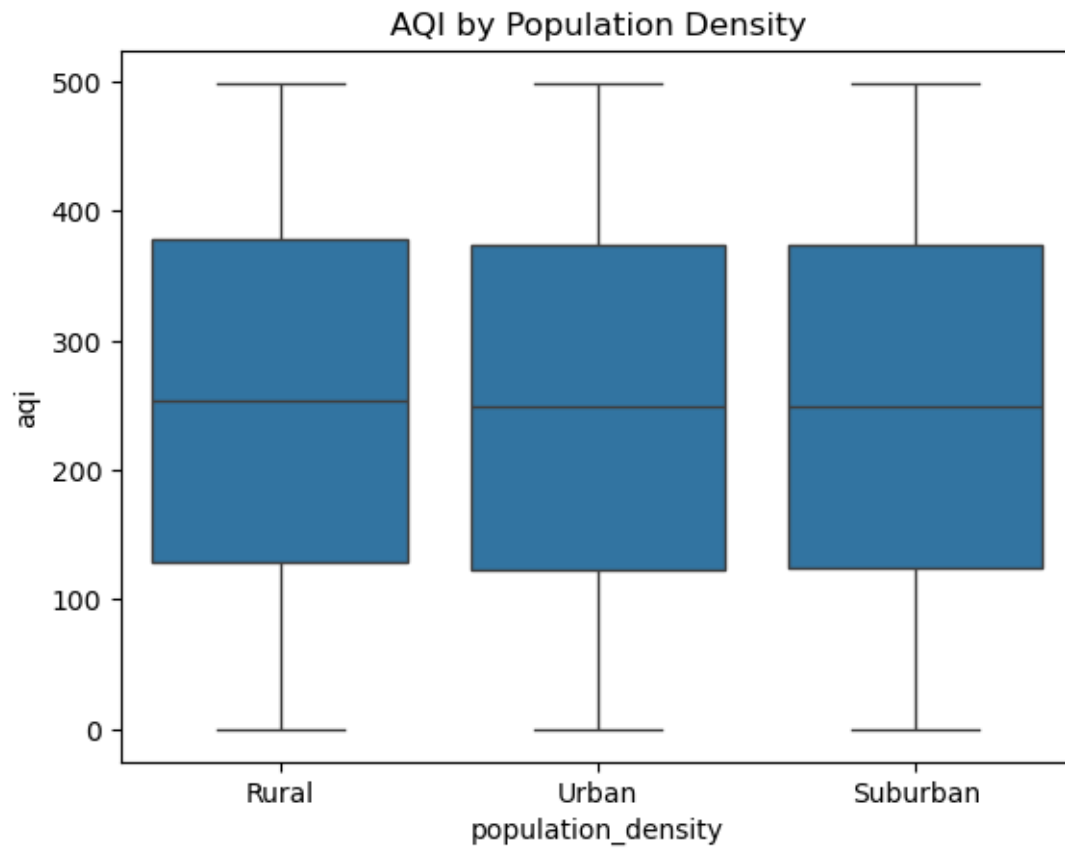
```
[41]: # Scatter plot showing the relationship between temperature and hospital_
      ↪ admissions
      sns.scatterplot(data=df, x='temperature', y='hospital_admissions')
      plt.title('Temperature vs Admissions')
      plt.show()
```



```
[43]: # Average Air Quality Index (AQI) by population density
      df.groupby('population_density')['aqi'].mean()
```

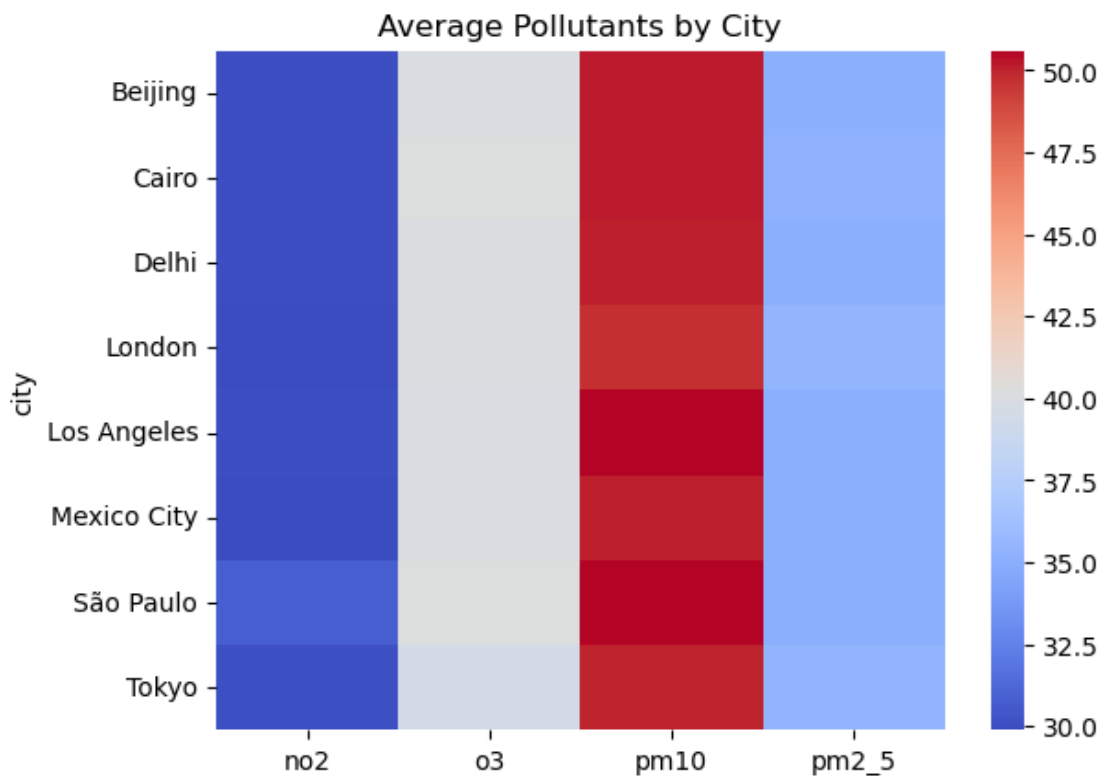
```
[43]: population_density
      Rural      252.397079
      Suburban  248.999131
      Urban    249.052207
      Name: aqi, dtype: float64
```

```
[45]: # Boxplot of AQI levels by population density
sns.boxplot(data=df, x='population_density', y='aqi')
plt.title('AQI by Population Density')
plt.show()
```



```
[49]: # Heatmap showing average pollution levels (PM2.5, PM10, NO2, O3) by city

pollutants = df.pivot_table(index='city', values=['pm2_5','pm10','no2','o3'],
    ↪aggfunc='mean')
sns.heatmap(pollutants, cmap='coolwarm')
plt.title('Average Pollutants by City')
plt.show()
```



```
[61]: # Top cities with the highest total hospital admissions
df.groupby('city')['hospital_admissions'].sum().nlargest(10)
```

```
[61]: city
Delhi      212893
Beijing    178631
Mexico City 107117
Los Angeles 72215
London     56281
Tokyo      49351
Cairo      21893
```

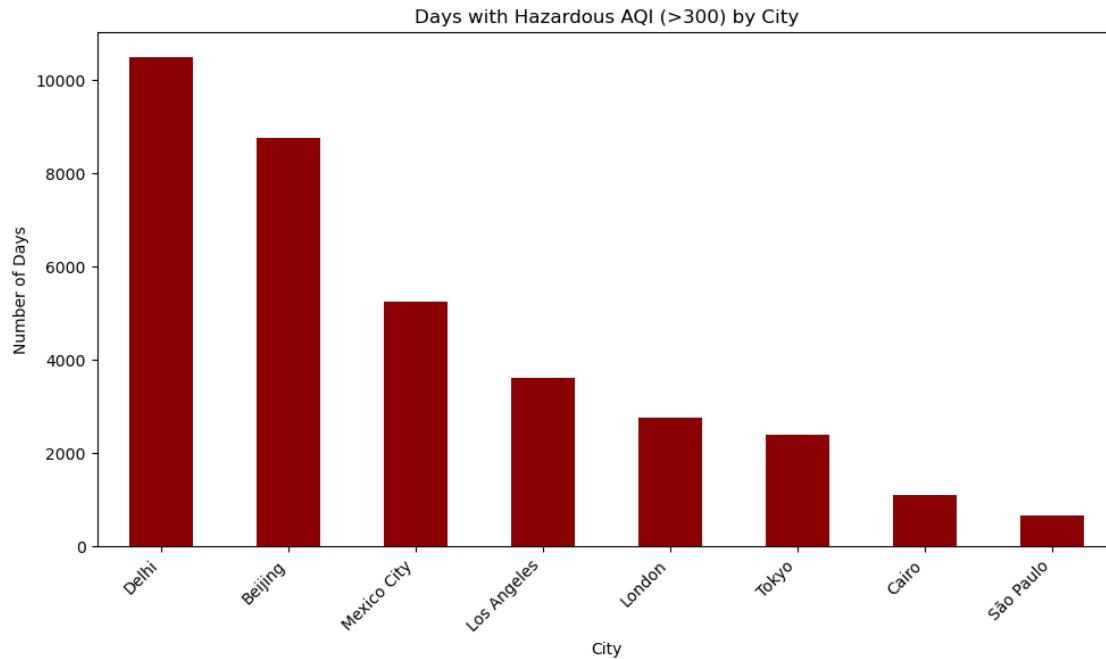
São Paulo 13901
Name: hospital_admissions, dtype: int64

```
[105]: #City Comparison of Extreme Pollution Days  
# Count days with AQI > 300 (Hazardous) by city  
  
hazardous_days = df[df['aqi'] > 300].groupby('city').size().  
    ↪sort_values(ascending=False)  
print("Days with Hazardous AQI (>300) by City:")  
print(hazardous_days)  
plt.figure(figsize=(10, 6))  
hazardous_days.plot(kind='bar', color='darkred')  
plt.title('Days with Hazardous AQI (>300) by City')  
plt.xlabel('City')  
plt.ylabel('Number of Days')  
plt.xticks(rotation=45, ha='right')  
plt.tight_layout()  
plt.show()
```

Days with Hazardous AQI (>300) by City:

city	
Delhi	10492
Beijing	8752
Mexico City	5250
Los Angeles	3617
London	2771
Tokyo	2402
Cairo	1101
São Paulo	676

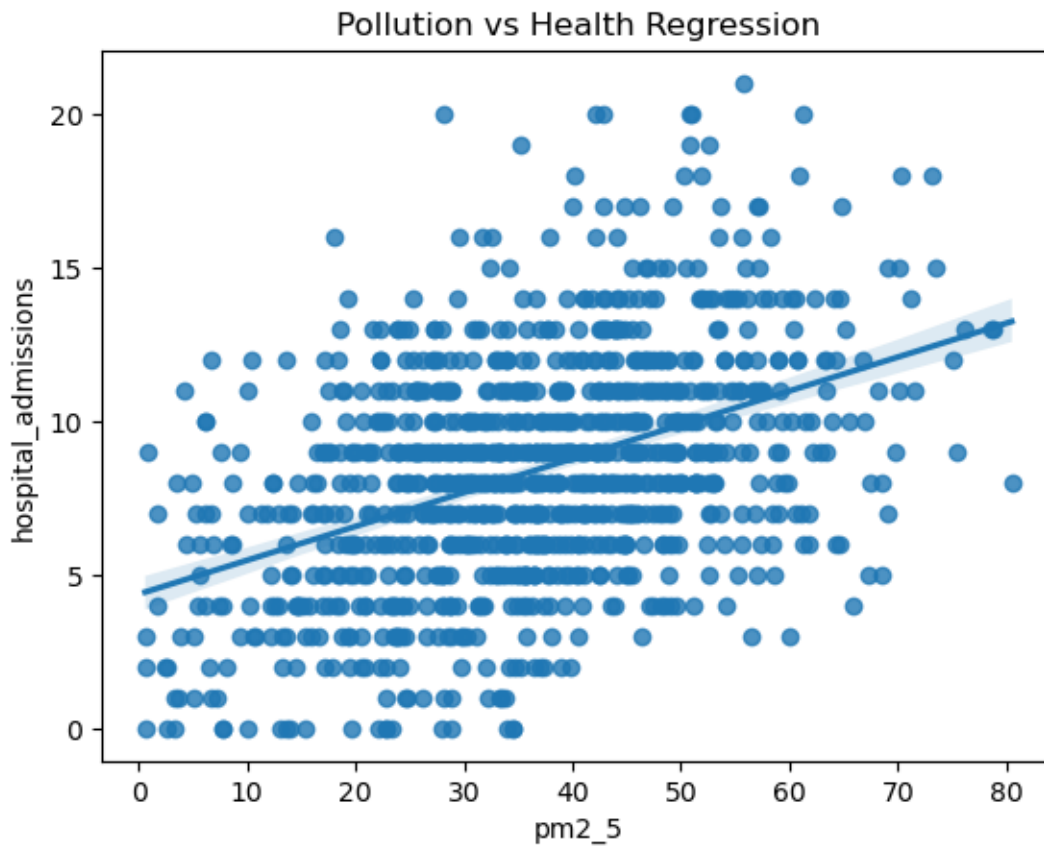
dtype: int64



```
[71]: # Annual average AQI over the years
mask = (df['date']>='2020-01-01')&(df['date']<'2021-01-01')
df[mask].groupby('city')['aqi'].mean()
```

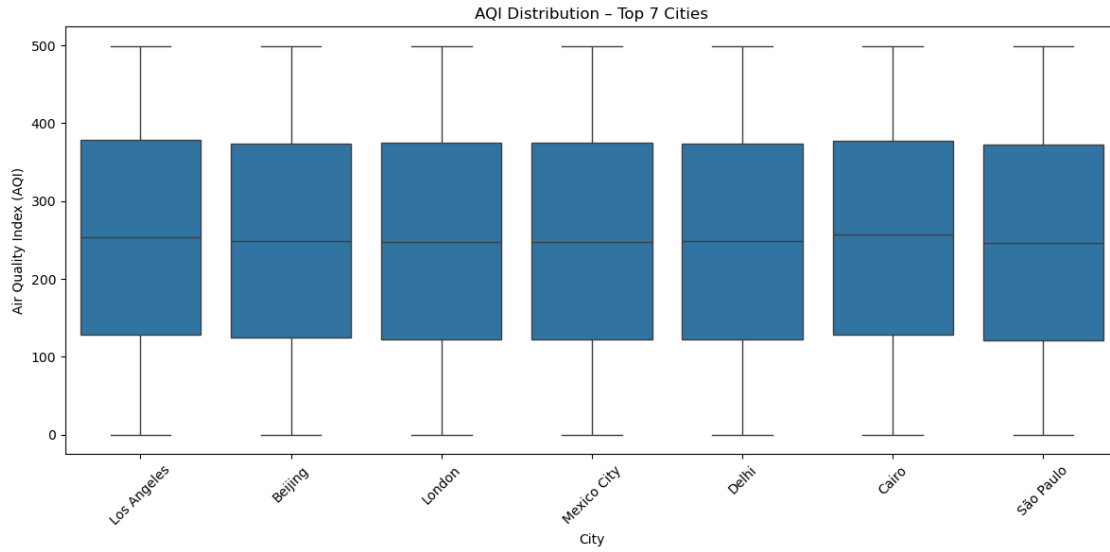
```
[71]: city
Beijing      273.400000
Cairo        236.916667
Delhi        245.412281
London       221.064516
Los Angeles  268.189189
Mexico City  249.177419
São Paulo    392.000000
Tokyo        237.458333
Name: aqi, dtype: float64
```

```
[85]: #Simple Linear Regression Analysis: Relationship between PM2.5 and Hospital
      ↳ Admissions
      sns.regplot(x='pm2_5', y='hospital_admissions', data=df.sample(1000))
      plt.title('Pollution vs Health Regression')
      plt.show()
```

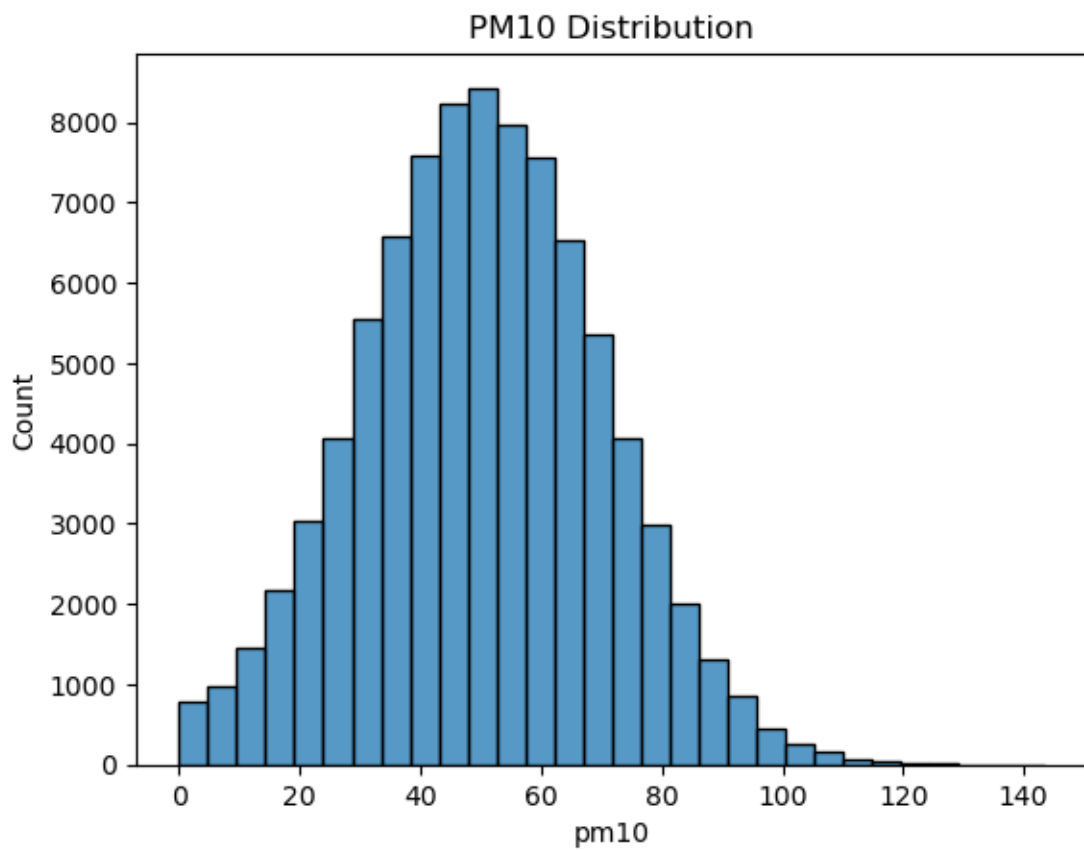


```
[95]: # 18. AQI Distribution for the Top 7 Most Polluted Cities
      top7 = df.groupby('city')['aqi'].mean().nlargest(7).index

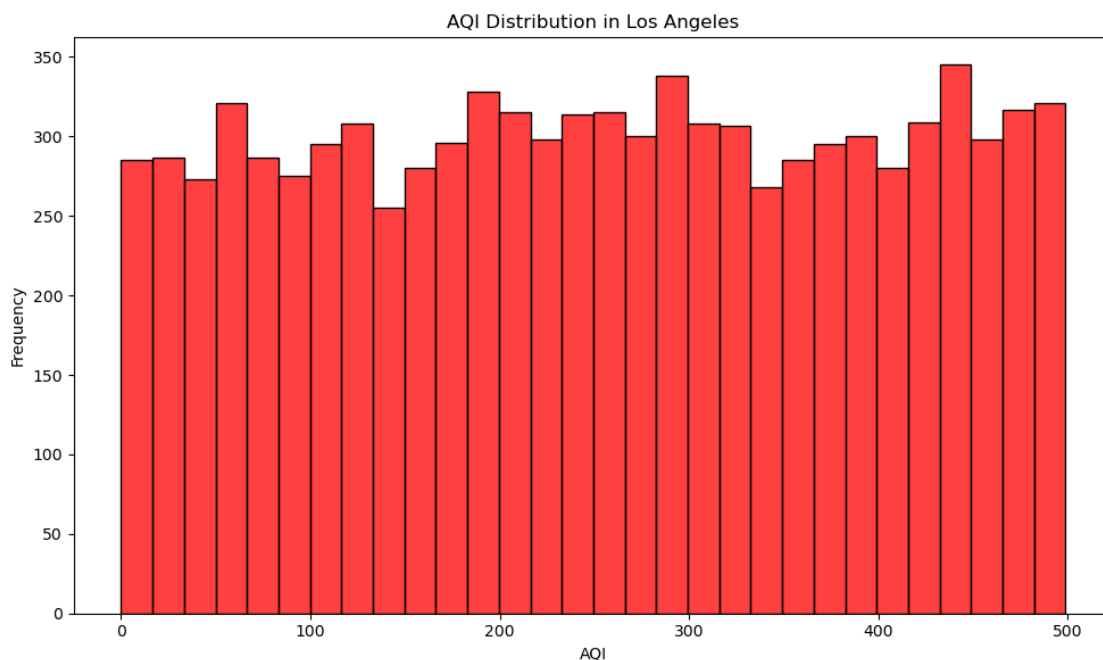
      plt.figure(figsize=(12, 6))
      sns.boxplot(data=df[df['city'].isin(top7)], x='city', y='aqi')
      plt.title('AQI Distribution - Top 7 Cities')
      plt.xticks(rotation=45)
      plt.ylabel('Air Quality Index (AQI)')
      plt.xlabel('City')
      plt.tight_layout()
      plt.show()
```



```
[93]: sns.histplot(df['pm10'], bins=30)
plt.title('PM10 Distribution')
plt.show()
```



```
[103]: # City-Specific AQI Histogram
# Plot AQI distribution for top polluted city
top_city = df.groupby('city')['aqi'].mean().idxmax()
plt.figure(figsize=(10, 6))
sns.histplot(df[df['city'] == top_city]['aqi'], bins=30, color='red')
plt.title(f'AQI Distribution in {top_city}')
plt.xlabel('AQI')
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()
```



```
[99]: # Weekend vs. Weekday AQI Comparison
# Compare AQI on weekends vs. weekdays
df['day_of_week'] = df['date'].dt.dayofweek
df['is_weekend'] = df['day_of_week'].isin([5, 6])
aqi_weekend = df.groupby('is_weekend')['aqi'].mean()
print("Average AQI: Weekday vs. Weekend:")
print(aqi_weekend)
plt.figure(figsize=(8, 6))
aqi_weekend.plot(kind='bar', color=['blue', 'green'])
plt.title('Average AQI: Weekday vs. Weekend')
plt.xlabel('Day Type')
```

```
plt.ylabel('Average AQI')
plt.xticks([0, 1], ['Weekday', 'Weekend'], rotation=0)
plt.tight_layout()
plt.show()
```

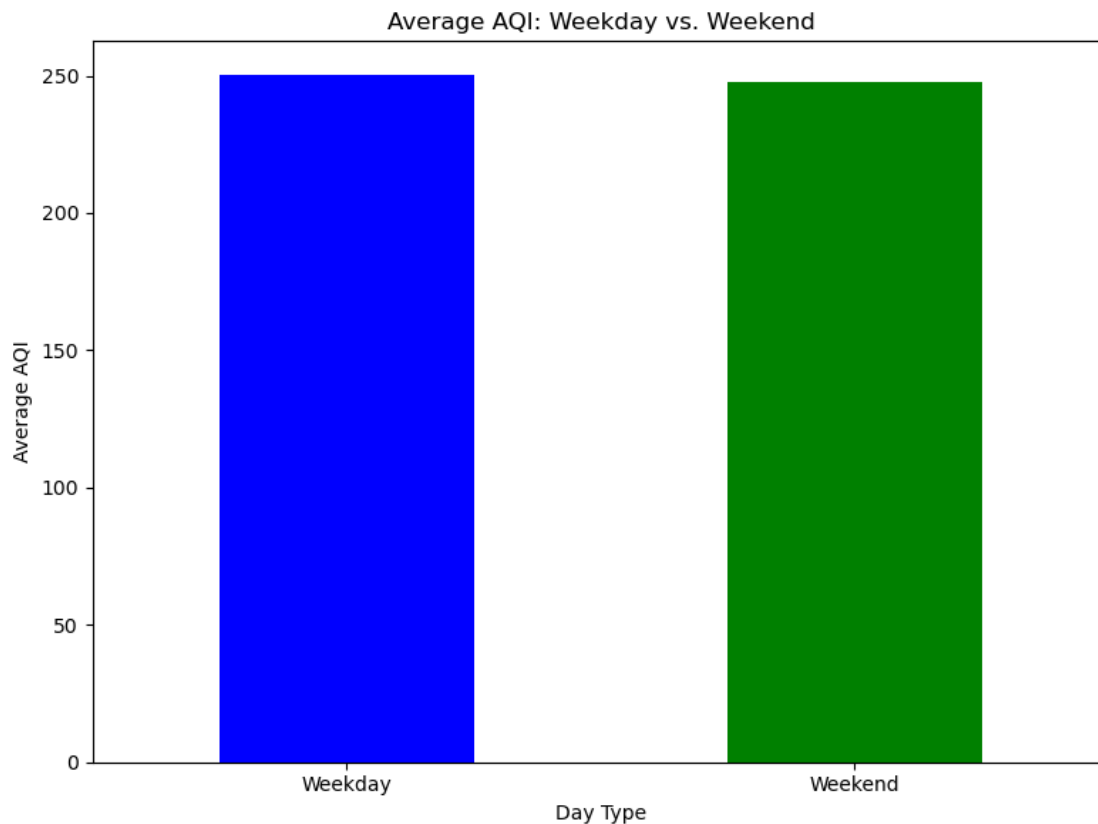
Average AQI: Weekday vs. Weekend:

is_weekend

False 250.043002

True 247.688078

Name: aqi, dtype: float64



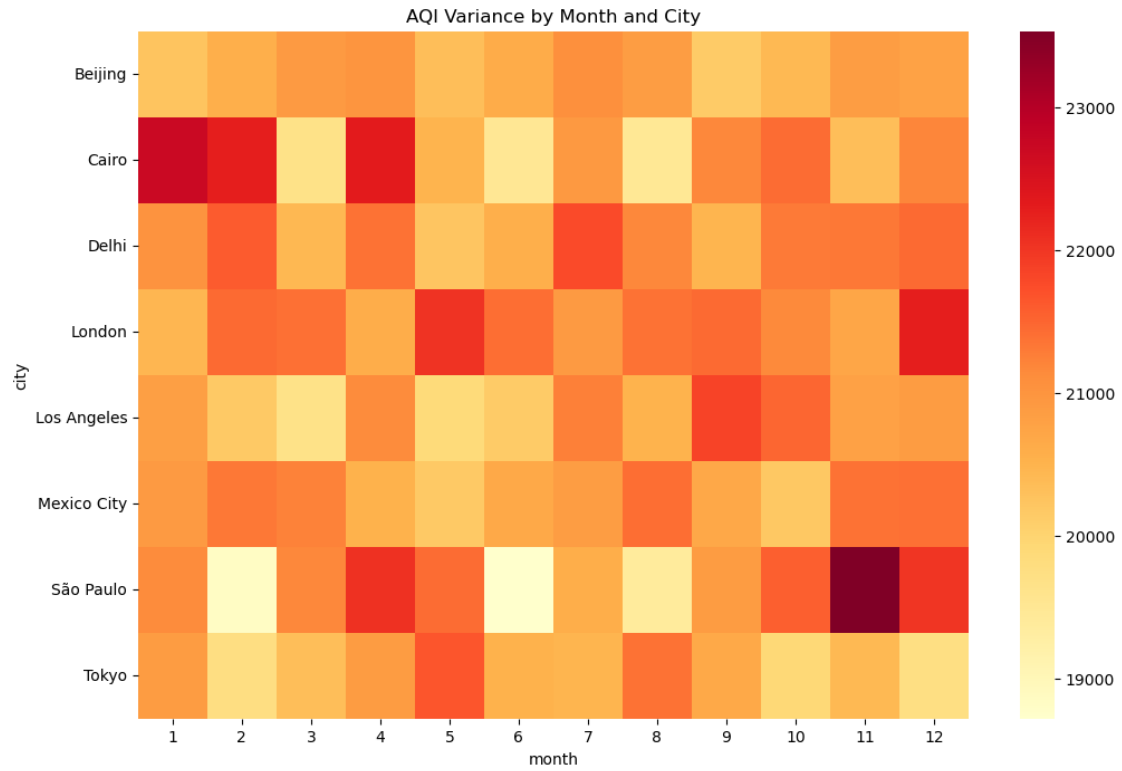
```
[109]: #Monthly AQI Variance by City
# Calculate variance of AQI by month for each city
df['month'] = df['date'].dt.month
aqi_variance = df.groupby(['city', 'month'])['aqi'].var().unstack()
print("AQI Variance by Month and City:")
print(aqi_variance)
plt.figure(figsize=(12, 8))
sns.heatmap(aqi_variance, cmap='YlOrRd')
plt.title('AQI Variance by Month and City')
plt.show()
```

AQI Variance by Month and City:

month	1	2	3	4	\
city					
Beijing	20251.681782	20576.396850	20909.597959	21001.512997	
Cairo	22705.849318	22288.054296	19675.014705	22323.321760	
Delhi	21026.724436	21608.593142	20438.884101	21371.467676	
London	20459.647485	21444.243397	21388.372484	20603.003784	
Los Angeles	20834.124699	20175.640740	19664.042579	21125.159225	
Mexico City	20906.085357	21322.747489	21224.460180	20530.157560	
São Paulo	21133.356530	18837.132799	21166.645705	22055.874986	
Tokyo	20889.843681	19752.248083	20345.702500	20886.261658	

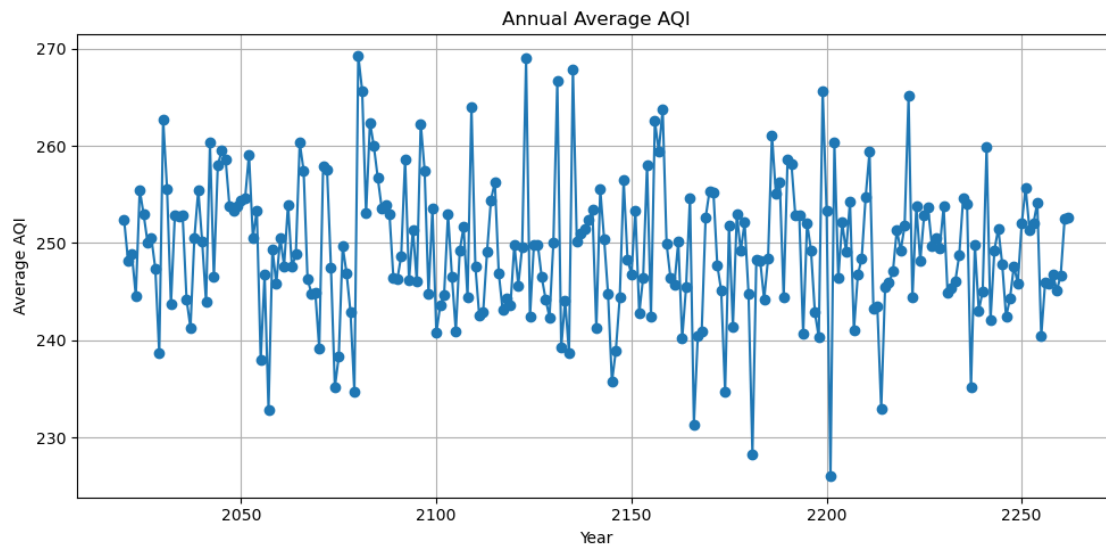
month	5	6	7	8	\
city					
Beijing	20349.831404	20626.295799	21085.438105	20878.264081	
Cairo	20485.494110	19504.547400	20930.691502	19477.923988	
Delhi	20235.484214	20564.806929	21780.744638	21177.805578	
London	22037.544638	21421.969065	20908.783240	21374.595598	
Los Angeles	19881.123054	20140.608000	21245.291729	20514.335777	
Mexico City	20175.782638	20670.814945	20870.601782	21409.540250	
São Paulo	21425.295260	18717.854806	20586.918321	19386.432206	
Tokyo	21668.044243	20540.277919	20484.143195	21378.548921	

month	9	10	11	12	
city					
Beijing	20140.346190	20424.515405	20870.025037	20775.460928	
Cairo	21169.391964	21433.884077	20334.508889	21192.104096	
Delhi	20481.948432	21308.092310	21318.203511	21453.581290	
London	21444.010012	21158.324686	20722.905136	22270.714671	
Los Angeles	21848.777871	21487.742809	20803.065800	20896.395207	
Mexico City	20688.794599	20184.762181	21369.139185	21387.439433	
São Paulo	20895.014631	21557.394465	23530.107178	22013.838936	
Tokyo	20659.706269	19910.498150	20412.535430	19739.822663	



```
[67]: # Annual average AQI over the years
annual_avg_aqi = df.groupby(df['date'].dt.year)['aqi'].mean()

plt.figure(figsize=(10, 5))
annual_avg_aqi.plot(marker='o')
plt.title('Annual Average AQI')
plt.xlabel('Year')
plt.ylabel('Average AQI')
plt.grid(True)
plt.tight_layout()
plt.show()
```




```
[73]: sns.  
      ↪heatmap(df[['aqi','pm2_5','pm10','no2','o3','temperature','hospital_admissions']]).  
      ↪corr(), annot=True)  
plt.title('Correlation Matrix')  
plt.show()
```