

# Synthetic HR Burnout Dataset

July 7, 2025

```
[125]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.cluster import KMeans
```

```
[126]: df = pd.read_csv('../data/synthetic_employee_burnout.csv')
print(df.head())
print(df.info())
```

	Name	Age	Gender	JobRole	Experience	WorkHoursPerWeek	\
0	Max Ivanov	32	Male	Analyst	3	60	
1	Max Wang	40	Female	Engineer	9	47	
2	Nina Petrov	33	Female	Engineer	2	44	
3	John Ivanov	35	Female	Manager	6	44	
4	John Wang	59	Male	Sales	8	38	

	RemoteRatio	SatisfactionLevel	StressLevel	Burnout
0	21	4.40	1	0
1	67	2.09	2	0
2	20	2.58	3	0
3	70	3.23	8	0
4	46	4.41	1	0

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 2000 entries, 0 to 1999

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	Name	2000 non-null	object
1	Age	2000 non-null	int64
2	Gender	2000 non-null	object
3	JobRole	2000 non-null	object
4	Experience	2000 non-null	int64
5	WorkHoursPerWeek	2000 non-null	int64
6	RemoteRatio	2000 non-null	int64
7	SatisfactionLevel	2000 non-null	float64
8	StressLevel	2000 non-null	int64

```
9    Burnout                2000 non-null    int64
dtypes: float64(1), int64(6), object(3)
memory usage: 156.4+ KB
None
```

```
[127]: # Missing values are checked for each column
print("Missing values:\n", df.isna().sum())
```

```
Missing values:
Name                0
Age                 0
Gender              0
JobRole             0
Experience           0
WorkHoursPerWeek    0
RemoteRatio         0
SatisfactionLevel   0
StressLevel         0
Burnout             0
dtype: int64
```

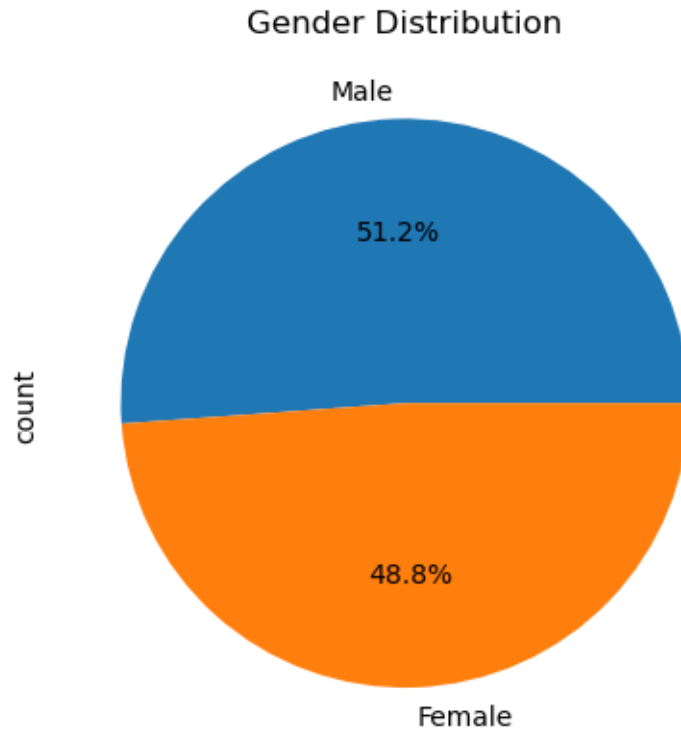
```
[128]: # Average age by burnout status
# This shows how age may relate to burnout levels in employees
avg_age_by_burnout = df.groupby("Burnout")["Age"].mean()
print(avg_age_by_burnout)
```

```
Burnout
0    40.682523
1    40.868217
Name: Age, dtype: float64
```

```
[129]: # Number of individuals by gender
# Gives a distribution of gender in the dataset
gender_count = df["Gender"].value_counts()
print(gender_count)
```

```
Gender
Male    1023
Female   977
Name: count, dtype: int64
```

```
[130]: df["Gender"].value_counts().plot(kind="pie", autopct='%1.1f%%')
plt.title("Gender Distribution")
plt.show()
```



```
[131]: # Average work hours per job role
# Helps identify roles with heavier work schedules
avg_hours_by_role = df.groupby("JobRole")["WorkHoursPerWeek"].mean()
print(avg_hours_by_role)
```

```
JobRole
Analyst      49.452785
Engineer     49.971503
HR           49.741688
Manager      49.310263
Sales        49.496164
Name: WorkHoursPerWeek, dtype: float64
```

```
[132]: # Burnout rate by gender
# Measures how burnout levels vary across genders
burnout_by_gender = df.groupby("Gender")["Burnout"].mean()
print(burnout_by_gender)
```

```
Gender
Female    0.062436
Male      0.066471
Name: Burnout, dtype: float64
```

```
[133]: # Top 3 most stressful job roles
# Identifies roles with the highest average stress level
top_stress_roles = df.groupby("JobRole")["StressLevel"].mean().nlargest(3)
print(top_stress_roles)
```

```
JobRole
Sales      5.583120
HR         5.524297
Manager    5.470167
Name: StressLevel, dtype: float64
```

```
[134]: # Distribution of remote work ratio
# Shows how common different levels of remote work are
remote_ratio_dist = df["RemoteRatio"].value_counts()
print(remote_ratio_dist.head())
```

```
RemoteRatio
22    32
42    30
82    28
53    28
45    28
Name: count, dtype: int64
```

```
[135]: # Average satisfaction by burnout level
# Analyzes how satisfaction changes with burnout
satisfaction_by_burnout = df.groupby("Burnout")["SatisfactionLevel"].mean()
print(satisfaction_by_burnout)
```

```
Burnout
0    3.065783
1    1.971938
Name: SatisfactionLevel, dtype: float64
```

```
[136]: # Average experience by job role
# Determines experience levels across different job roles
exp_by_role = df.groupby("JobRole")["Experience"].mean()
print(exp_by_role)
```

```
JobRole
Analyst      10.266344
Engineer     9.898964
HR           10.053708
Manager      10.093079
Sales        10.046036
Name: Experience, dtype: float64
```

```
[137]: # Visual comparison of work hours between burned out and non-burned out employees
df.boxplot(column="WorkHoursPerWeek", by="Burnout")
```


```
plt.title("Work Hours by Burnout Status")
plt.suptitle("")
plt.show()
```



```
[138]: # Top 10 individuals with highest stress
# Highlights names with highest average stress levels
top_stress_names = df.groupby("Name")["StressLevel"].mean().nlargest(10)
print(top_stress_names)
```

```
Name
Dina Ivanov    7.090909
Leo Garcia     6.933333
John Lee       6.888889
Kate Smith     6.666667
Ivan Smith     6.500000
Kate Wang      6.500000
Kate Garcia    6.444444
Dina Brown     6.411765
Max Brown      6.388889
Alex Kim       6.347826
Name: StressLevel, dtype: float64
```

```
[139]: plt.figure(figsize=(10, 6))
top_stress_names.plot(kind="bar", color="red")
plt.title("Top 10 Individuals with Highest Average Stress Levels")
plt.xlabel("Name")
plt.ylabel("Average Stress Level")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()
plt.grid(axis='y')
plt.show()
```



```
[140]: # Burnout rate by job role
# Shows average burnout rate for each job category
burnout_by_role = df.groupby("JobRole")["Burnout"].mean()
print(burnout_by_role)
```

```
JobRole
Analyst      0.048426
Engineer     0.064767
HR           0.066496
Manager      0.066826
Sales        0.076726
Name: Burnout, dtype: float64
```

```
[141]: # Average age by gender
# Helps identify if age distribution varies by gender
avg_age_by_gender = df.groupby("Gender")["Age"].mean()
print(avg_age_by_gender)
```

```
Gender
Female    40.580348
Male      40.803519
Name: Age, dtype: float64
```

```
[142]: # Satisfaction level by experience
# Shows how employee satisfaction changes with years of experience
satisfaction_by_exp = df.groupby("Experience")["SatisfactionLevel"].mean()
print(satisfaction_by_exp.head())
```

```
Experience
0    3.023735
1    3.096416
2    2.892133
3    3.082984
4    3.112479
Name: SatisfactionLevel, dtype: float64
```

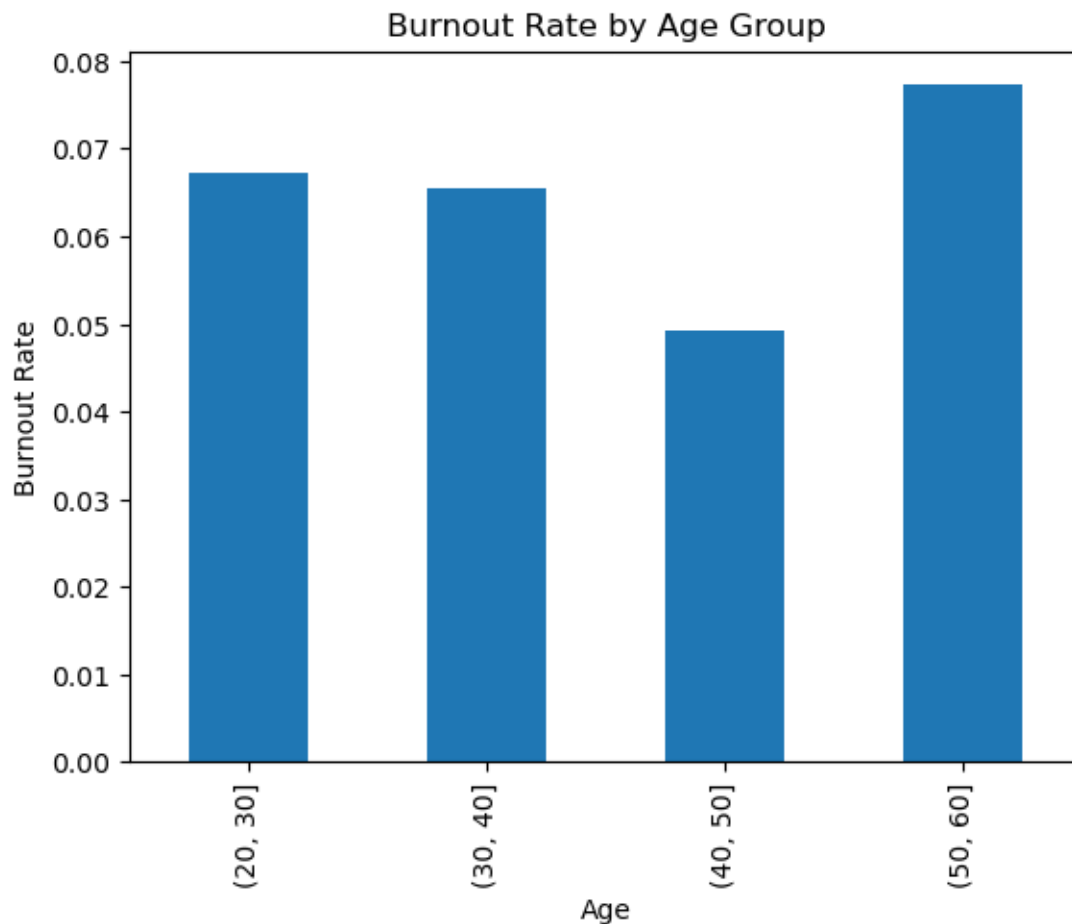
```
[143]: # Number of people per age group
# Breaks down population into age bins to analyze age distribution
age_bins = pd.cut(df["Age"], bins=[20, 30, 40, 50, 60])

age_dist = df.groupby(age_bins, observed=True)["Name"].count()
print(age_dist)
```

```
Age
(20, 30]    490
(30, 40]    504
(40, 50]    527
(50, 60]    479
Name: Name, dtype: int64
```

```
[144]: age_bins = pd.cut(df["Age"], bins=[20, 30, 40, 50, 60])

df.groupby(age_bins, observed=True)["Burnout"].mean().plot(kind="bar")
plt.title("Burnout Rate by Age Group")
plt.ylabel("Burnout Rate")
plt.show()
```



```
[145]: # This shows the average burnout level for each job role, helping identify
        ↪ high-risk positions.
burnout_by_role = df.groupby("JobRole")["Burnout"].mean()
print(burnout_by_role)
```

```
JobRole
Analyst      0.048426
Engineer     0.064767
HR           0.066496
Manager      0.066826
Sales        0.076726
```



Name: Burnout, dtype: float64

```
[146]: # Identifies which roles are most suitable for remote work based on their  
       ↪ average remote ratio.
```

```
top_remote_roles = df.groupby("JobRole")["RemoteRatio"].mean().nlargest(5)  
print(top_remote_roles)
```

```
JobRole  
Engineer    51.538860  
HR          50.744246  
Analyst     49.704600  
Sales       49.682864  
Manager     48.346062
```

Name: RemoteRatio, dtype: float64

```
[147]: # Compares workload between genders to ensure equitable work distribution.  
hours_by_gender = df.groupby("Gender")["WorkHoursPerWeek"].mean()  
print(hours_by_gender)
```

```
Gender  
Female     49.518936  
Male       49.653959
```

Name: WorkHoursPerWeek, dtype: float64

```
[148]: # Top 10 individuals with the highest average satisfaction level  
top_satisfaction_names = df.groupby("Name")["SatisfactionLevel"].mean().  
       ↪ nlargest(10)  
print(top_satisfaction_names)
```

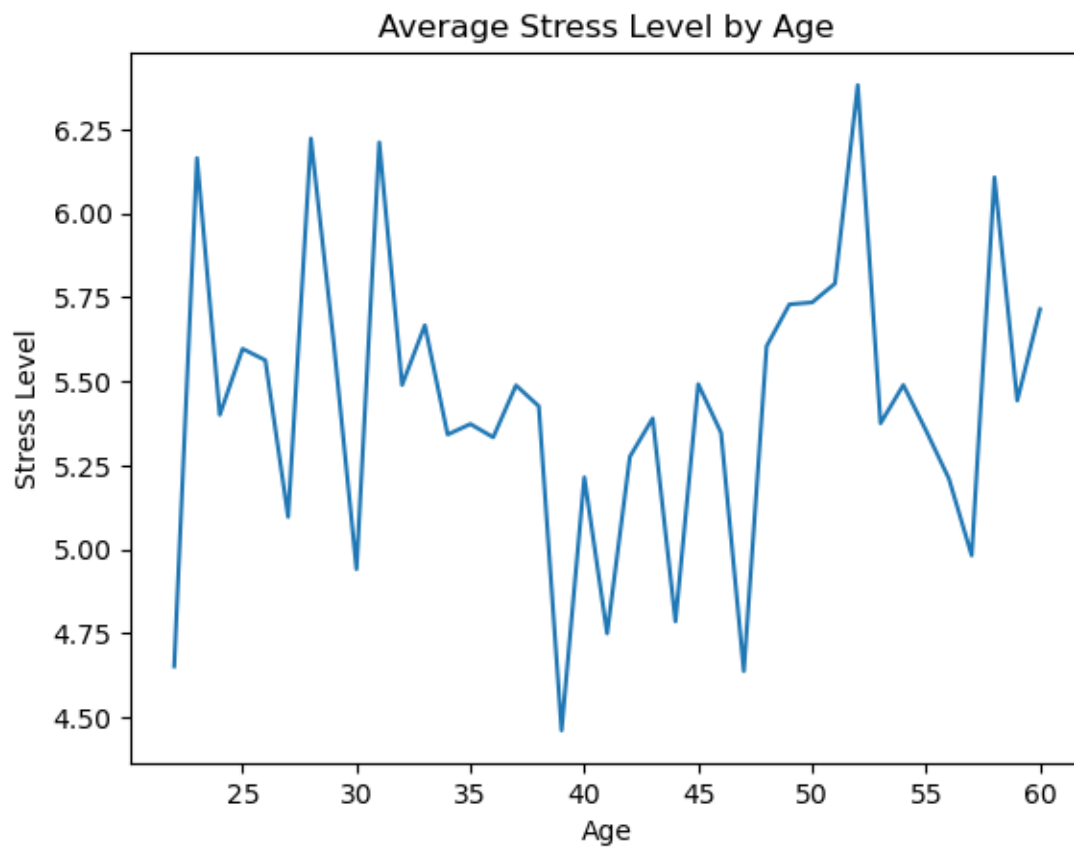
```
Name  
Alex Brown    3.688125  
John Smith    3.632857  
Sam Chen      3.492667  
Lily Brown    3.473125  
Lily Ivanov   3.385909  
Kate Garcia   3.374444  
Lily Wang     3.357813  
Alex Ivanov   3.351579  
Max Kim       3.349412  
Dina Wang     3.348421  
Name: SatisfactionLevel, dtype: float64
```

```
[149]: burnout_by_stress = df.groupby("StressLevel")["Burnout"].mean()  
print(burnout_by_stress)
```

```
StressLevel  
1    0.000000  
2    0.000000  
3    0.000000  
4    0.000000
```

```
5    0.000000
6    0.000000
7    0.000000
8    0.211538
9    0.219626
10   0.215909
Name: Burnout, dtype: float64
```

```
[150]: # Visualizes how stress levels vary with age - can indicate support needs for
↳ specific age groups.
df.groupby("Age")["StressLevel"].mean().plot(kind="line")
plt.title("Average Stress Level by Age")
plt.ylabel("Stress Level")
plt.show()
```



```
[151]: pivot = df.pivot_table(values="StressLevel", index="SatisfactionLevel",
    ↪aggfunc="mean")
sns.heatmap(pivot)
plt.title("Stress vs Satisfaction")
plt.show()
```



```
[152]: # Create a heatmap to show how stress levels vary with work hours
pivot = df.pivot_table(values="StressLevel", index="WorkHoursPerWeek",
↳aggfunc="mean")
sns.heatmap(pivot)
plt.title("Stress Level by Work Hours")
plt.show()
```



```
[153]: # Group employees into clusters using KMeans based on their stress and
↳satisfaction levels
X_cluster = df[["StressLevel", "SatisfactionLevel"]]
kmeans = KMeans(n_clusters=3, random_state=42).fit(X_cluster)
df["Cluster"] = kmeans.labels_
print(df.groupby("Cluster")[["StressLevel", "SatisfactionLevel"]].mean())
```

	StressLevel	SatisfactionLevel
Cluster		
0	5.483418	2.999388
1	1.966019	2.951764
2	8.946488	3.034699

```
[154]: # Compare burnout rates by experience levels and gender
satisfaction_by_role_gender = df.groupby(["JobRole",
↳ "Gender"])["SatisfactionLevel"].mean()
print(satisfaction_by_role_gender)
```

JobRole	Gender	
Analyst	Female	2.862033
	Male	3.023420
Engineer	Female	3.005677
	Male	2.945103
HR	Female	3.023143
	Male	2.975028
Manager	Female	3.042133
	Male	2.995096
Sales	Female	3.018901
	Male	3.038612

Name: SatisfactionLevel, dtype: float64

```
[155]: #Basic Predictive Modeling for Burnout (Logistic Regression)
# Predict burnout using logistic regression based on key features

X = df[["Age", "WorkHoursPerWeek", "RemoteRatio", "SatisfactionLevel",
↳ "StressLevel"]]
y = df["Burnout"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
model = LogisticRegression().fit(X_train, y_train)
print("Accuracy:", model.score(X_test, y_test))
```

Accuracy: 0.95

```
[156]: df.groupby("JobRole")["StressLevel"].mean().plot(kind="bar")
plt.title("Average Stress Level by Job Role")
plt.ylabel("Stress Level")
plt.show()
```

