

Ecommerce Consumer Behavior Insights (ECBI)

July 13, 2025

```
[73]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[74]: df = pd.read_csv(r'../data/Ecommerce_Consumer_Behavior_Analysis_Data.csv')

print(df.head())
print(df.info())
```

	Customer_ID	Age	Gender	Income_Level	Marital_Status	Education_Level	\
0	37-611-6911	22	Female	Middle	Married	Bachelor's	
1	29-392-9296	49	Male	High	Married	High School	
2	84-649-5117	24	Female	Middle	Single	Master's	
3	48-980-6078	29	Female	Middle	Single	Master's	
4	91-170-9072	33	Female	Middle	Widowed	High School	

	Occupation	Location	Purchase_Category	Purchase_Amount	...	\
0	Middle	Évry	Gardening & Outdoors	\$333.80	...	
1	High	Huocheng	Food & Beverages	\$222.22	...	
2	High	Huzhen	Office Supplies	\$426.22	...	
3	Middle	Wiwilí	Home Appliances	\$101.31	...	
4	Middle	Nara	Furniture	\$211.70	...	

	Customer_Satisfaction	Engagement_with_Ads	Device_Used_for_Shopping	\
0	7	NaN	Tablet	
1	5	High	Tablet	
2	7	Low	Smartphone	
3	1	NaN	Smartphone	
4	10	NaN	Smartphone	

	Payment_Method	Time_of_Purchase	Discount_Used	\
0	Credit Card	3/1/2024	True	
1	PayPal	4/16/2024	True	
2	Debit Card	3/15/2024	True	
3	Other	10/4/2024	True	
4	Debit Card	1/30/2024	False	

	Customer_Loyalty_Program_Member	Purchase_Intent	Shipping_Preference	\
--	---------------------------------	-----------------	---------------------	---

0	False	Need-based	No Preference
1	False	Wants-based	Standard
2	True	Impulsive	No Preference
3	True	Need-based	Express
4	False	Wants-based	No Preference

Time_to_Decision	
0	2
1	6
2	3
3	10
4	4

```
[5 rows x 28 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 28 columns):
```

#	Column	Non-Null Count	Dtype
0	Customer_ID	1000 non-null	object
1	Age	1000 non-null	int64
2	Gender	1000 non-null	object
3	Income_Level	1000 non-null	object
4	Marital_Status	1000 non-null	object
5	Education_Level	1000 non-null	object
6	Occupation	1000 non-null	object
7	Location	1000 non-null	object
8	Purchase_Category	1000 non-null	object
9	Purchase_Amount	1000 non-null	object
10	Frequency_of_Purchase	1000 non-null	int64
11	Purchase_Channel	1000 non-null	object
12	Brand_Loyalty	1000 non-null	int64
13	Product_Rating	1000 non-null	int64
14	Time_Spent_on_Product_Research(hours)	1000 non-null	float64
15	Social_Media_Influence	753 non-null	object
16	Discount_Sensitivity	1000 non-null	object
17	Return_Rate	1000 non-null	int64
18	Customer_Satisfaction	1000 non-null	int64
19	Engagement_with_Ads	744 non-null	object
20	Device_Used_for_Shopping	1000 non-null	object
21	Payment_Method	1000 non-null	object
22	Time_of_Purchase	1000 non-null	object
23	Discount_Used	1000 non-null	bool
24	Customer_Loyalty_Program_Member	1000 non-null	bool
25	Purchase_Intent	1000 non-null	object
26	Shipping_Preference	1000 non-null	object
27	Time_to_Decision	1000 non-null	int64

```
dtypes: bool(2), float64(1), int64(7), object(18)
```

memory usage: 205.2+ KB
None

```
[75]: df["Social_Media_Influence"] = df["Social_Media_Influence"].fillna("Unknown")  
df["Engagement_with_Ads"] = df["Engagement_with_Ads"].fillna("Unknown")
```

```
[76]: # Missing values are checked for each column  
print("Missing values:\n", df.isna().sum())
```

```
Missing values:  
Customer_ID          0  
Age                  0  
Gender               0  
Income_Level         0  
Marital_Status       0  
Education_Level      0  
Occupation           0  
Location             0  
Purchase_Category   0  
Purchase_Amount      0  
Frequency_of_Purchase 0  
Purchase_Channel     0  
Brand_Loyalty        0  
Product_Rating       0  
Time_Spent_on_Product_Research(hours) 0  
Social_Media_Influence 0  
Discount_Sensitivity 0  
Return_Rate          0  
Customer_Satisfaction 0  
Engagement_with_Ads  0  
Device_Used_for_Shopping 0  
Payment_Method       0  
Time_of_Purchase     0  
Discount_Used        0  
Customer_Loyalty_Program_Member 0  
Purchase_Intent      0  
Shipping_Preference  0  
Time_to_Decision     0  
dtype: int64
```

```
[77]: #the average age of customers for each income level  
avg_age_by_income = df.groupby("Income_Level")["Age"].mean()  
print(avg_age_by_income)  
avg_age_by_income = df.groupby("Income_Level")["Age"].mean()  
print(avg_age_by_income)
```

```
Income_Level  
High      34.231068  
Middle    34.381443
```

```
Name: Age, dtype: float64
Income_Level
High      34.231068
Middle    34.381443
Name: Age, dtype: float64
```

```
[78]: #the number of customers for each gender
gender_dist = df["Gender"].value_counts()
print(gender_dist)
```

```
Gender
Female      452
Male        449
Bigender     20
Agender      19
Genderfluid  17
Non-binary   16
Polygender   15
Genderqueer  12
Name: count, dtype: int64
```

```
[79]: #purchase amounts by removing '$' and converting to float
df['Purchase_Amount'] = df['Purchase_Amount'].replace('[\$,]', '', regex=True).
    ↪astype(float)
#the average purchase amount for each category
avg_amount_by_category = df.groupby("Purchase_Category")["Purchase_Amount"].
    ↪mean()
print(avg_amount_by_category.head())
```

```
Purchase_Category
Animal Feed      260.615909
Arts & Crafts    221.468235
Baby Products    272.500488
Beauty & Personal Care  233.676765
Books            300.613243
Name: Purchase_Amount, dtype: float64
```

```
[80]: #customers in each marital status category
marital_dist = df["Marital_Status"].value_counts()
print(marital_dist)
```

```
Marital_Status
Widowed      260
Married      253
Divorced     245
Single       242
Name: count, dtype: int64
```

```
[81]: #average satisfaction scores per channel
satisfaction_by_channel = df.
    ↳groupby("Purchase_Channel")["Customer_Satisfaction"].mean()
print(satisfaction_by_channel)
```

```
Purchase_Channel
In-Store      5.239264
Mixed         5.379412
Online        5.574850
Name: Customer_Satisfaction, dtype: float64
```

```
[82]: #total purchase amounts by device used
total_by_device = df.groupby("Device_Used_for_Shopping")["Purchase_Amount"].sum()
print(total_by_device)
```

```
Device_Used_for_Shopping
Desktop        93344.49
Smartphone     87717.59
Tablet         94001.80
Name: Purchase_Amount, dtype: float64
```


```
[83]: #the top 5 purchase categories with the highest total spending
top_categories = df.groupby("Purchase_Category")["Purchase_Amount"].sum().
    ↳nlargest(5)
print(top_categories)
```

```
Purchase_Category
Jewelry & Accessories    15139.36
Sports & Outdoors       14610.51
Electronics             13842.41
Software & Apps         13601.41
Toys & Games            13536.46
Name: Purchase_Amount, dtype: float64
```

```
[84]: #Compare average purchase amount between loyalty members vs non-members
loyalty_by_amount = df.
    ↳groupby("Customer_Loyalty_Program_Member")["Purchase_Amount"].mean()
print(loyalty_by_amount)
```

```
Customer_Loyalty_Program_Member
False      288.373026
True       261.266823
Name: Purchase_Amount, dtype: float64
```

```
[85]: df["Education_Level"].value_counts().plot(kind="bar")
plt.title("Distribution of Education Levels")
plt.ylabel("Count")
plt.show()
```



```
[86]: #average research time per age
research_by_age = df.groupby("Age")["Time_Spent_on_Product_Research(hours)"].
    ↪mean()
print(research_by_age.head())
```

```
Age
18    1.230000
19    0.921875
20    0.710526
21    1.109429
22    0.973684
```

Name: Time_Spent_on_Product_Research(hours), dtype: float64

```
[87]: #percentage of discount usage by gender
discount_by_gender = df.groupby("Gender")["Discount_Used"].
    ↪value_counts(normalize=True) * 100
print(discount_by_gender.head())
```

Gender	Discount_Used	
Agender	True	68.421053
	False	31.578947
Bigender	False	55.000000
	True	45.000000
Female	True	53.097345

Name: proportion, dtype: float64

```
[88]: #the average satisfaction scores for each income level
satisfaction_by_income = df.groupby("Income_Level")["Customer_Satisfaction"].
    ↪mean()
print(satisfaction_by_income)
```

Income_Level	
High	5.207767
Middle	5.602062

Name: Customer_Satisfaction, dtype: float64

```
[89]: #average time taken to make purchase decisions per category
decision_by_category = df.groupby("Purchase_Category")["Time_to_Decision"].mean()
print(decision_by_category.head())
```

Purchase_Category	
Animal Feed	7.818182
Arts & Crafts	8.058824
Baby Products	7.536585
Beauty & Personal Care	8.382353
Books	8.189189

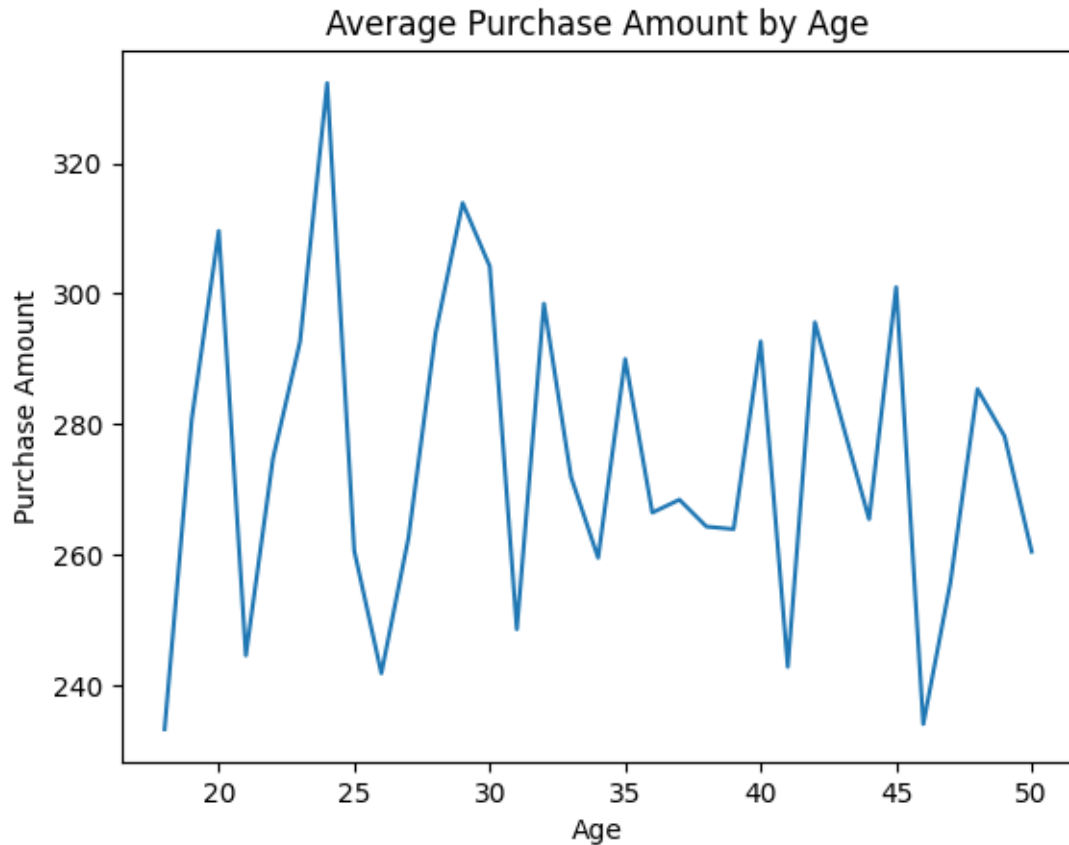
Name: Time_to_Decision, dtype: float64

```
[90]: # Group customers into age bins and count how many fall into each
age_bins = pd.cut(df["Age"], bins=[0, 30, 50, 70, 100])
age_dist = df.groupby(age_bins, observed=False)["Customer_ID"].count()
print(age_dist)
```

Age	
(0, 30]	387
(30, 50]	613
(50, 70]	0
(70, 100]	0

Name: Customer_ID, dtype: int64

```
[91]: df.groupby("Age")["Purchase_Amount"].mean().plot(kind="line")
plt.title("Average Purchase Amount by Age")
plt.ylabel("Purchase Amount")
plt.show()
```



```
[92]: #Compare average purchase amount based on ad engagement
ads_by_amount = df.groupby("Engagement_with_Ads")["Purchase_Amount"].mean()
print(ads_by_amount)
```


```
Engagement_with_Ads
High      277.658222
Low       268.259522
Medium    281.211107
Unknown   272.581875
Name: Purchase_Amount, dtype: float64
```



```
[93]: total_by_payment = df.groupby("Payment_Method")["Purchase_Amount"].sum()
print(total_by_payment)
```

```
Payment_Method
Cash          50072.47
Credit Card   52677.54
Debit Card     53552.04
Other          57797.50
PayPal         60964.33
Name: Purchase_Amount, dtype: float64
```

```
[94]: df["Payment_Method"].value_counts().plot(kind="pie", autopct='%1.1f%%')
plt.title("Distribution of Payment Methods")
plt.show()
```



```
[95]: #how purchase intent varies by gender
intent_by_gender = df.groupby("Gender")["Purchase_Intent"].value_counts()
print(intent_by_gender.head())
```

```
Gender    Purchase_Intent
Agender   Need-based      6
          Planned         5
          Wants-based     5
```

```
Impulsive      3
Bigender Need-based      7
Name: count, dtype: int64
```

```
[96]: #how social media influence affects purchase intent
influence_by_intent = df.groupby("Social_Media_Influence")["Purchase_Intent"].
    ↪value_counts()
print(influence_by_intent.head())
```

```
Social_Media_Influence Purchase_Intent
High
    Impulsive      73
    Wants-based    73
    Need-based     66
    Planned       56
Low
    Planned      72
Name: count, dtype: int64
```

```
[97]: #the number of customers per purchase channel
channel_dist = df["Purchase_Channel"].value_counts()
print(channel_dist)
```

```
Purchase_Channel
Mixed      340
Online     334
In-Store   326
Name: count, dtype: int64
```

```
[98]: #how satisfaction varies depending on purchase intent
satisfaction_by_intent = df.groupby("Purchase_Intent")["Customer_Satisfaction"].
    ↪mean()
print(satisfaction_by_intent)
```

```
Purchase_Intent
Impulsive      5.395161
Need-based     5.375000
Planned        5.230769
Wants-based    5.594378
Name: Customer_Satisfaction, dtype: float64
```

```
[99]: #average purchase amount with and without discount use
discount_by_amount = df.groupby("Discount_Used")["Purchase_Amount"].mean()
print(discount_by_amount)
```

```
Discount_Used
False      276.229436
True       273.992284
Name: Purchase_Amount, dtype: float64
```

```
[100]: #total spending by customers in each occupation  
total_by_occupation = df.groupby("Occupation")["Purchase_Amount"].sum()  
print(total_by_occupation)
```

```
Occupation  
High      140776.67  
Middle    134287.21  
Name: Purchase_Amount, dtype: float64
```

```
[101]: #the top 10 customers with the highest total purchase spending  
top_customers = df.groupby("Customer_ID")["Purchase_Amount"].sum().nlargest(10)  
print(top_customers)
```

```
Customer_ID  
60-470-3563    498.33  
13-848-5757    498.23  
15-663-7994    497.80  
86-257-9581    497.76  
15-421-1255    497.75  
72-830-1211    496.11  
85-467-6564    495.95  
72-590-6161    495.80  
69-394-1424    494.97  
59-261-4453    494.81  
Name: Purchase_Amount, dtype: float64
```

```
[102]: pivot = df.pivot_table(values="Purchase_Amount", index="Customer_Satisfaction",
    ↪aggfunc="mean")
sns.heatmap(pivot)
plt.title("Purchase Amount vs Customer Satisfaction")
plt.show()
```



```
[103]: #Identify the 10 locations with the highest total purchase amount
top_locations = df.groupby("Location")["Purchase_Amount"].sum().nlargest(10)
print(top_locations)
```

```
Location
Göteborg      1161.29
Oslo          1021.55
Punta Gorda    820.45
Magdalena      804.74
Hoolt          780.60
Veiga          779.89
San Carlos     722.39
Týn nad Vltavou 682.59
Cimara         673.61
Seleuš        672.99
```

Name: Purchase_Amount, dtype: float64

```
[104]: plt.figure(figsize=(10, 6))
top_locations.plot(kind="bar")
plt.title("Top 10 Locations by Total Purchase Amount")
plt.xlabel("Location")
plt.ylabel("Total Purchase Amount")
plt.xticks(rotation=45)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

